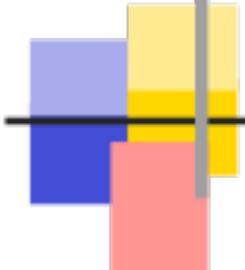


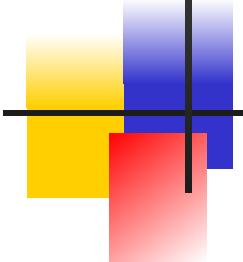
Lazy functional
programming for real



Tackling the Awkward Squad

Simon Peyton Jones

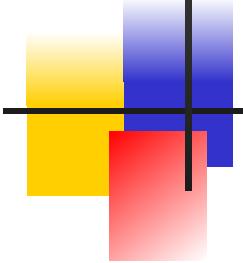
Microsoft Research



Beauty and the Beast

- Functional programming is beautiful, and many books tell us why
- But to write real applications, we must deal with un-beautiful issues “around the edges”:
 - Input/output
 - Concurrency
 - Error recovery
 - Foreign-language interfaces



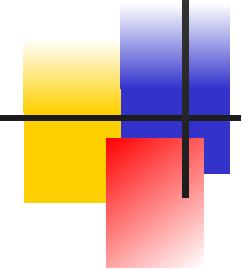


The direct approach

Do everything in “the usual way”

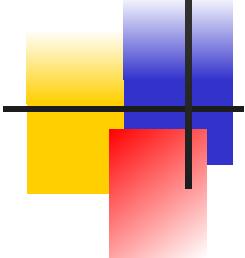
- I/O via “functions” with side effects
`putchar 'x'` + `putchar 'y'`
- Concurrency via operating system threads; OS calls mapped to “functions”
- Error recovery via exceptions
- Foreign language *procedures* mapped to “functions”

The lazy hair shirt



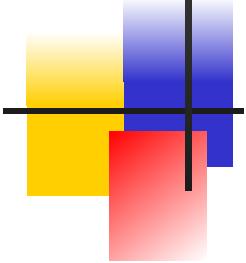
In a lazy functional language, like Haskell, order of evaluation is deliberately undefined.

- **`putchar 'x' + putchar 'y'`**
Output depends on evaluation order of (+)
- **`[putchar 'x', putchar 'y']`**
Output (if any) depends on how the consumer evaluates the list



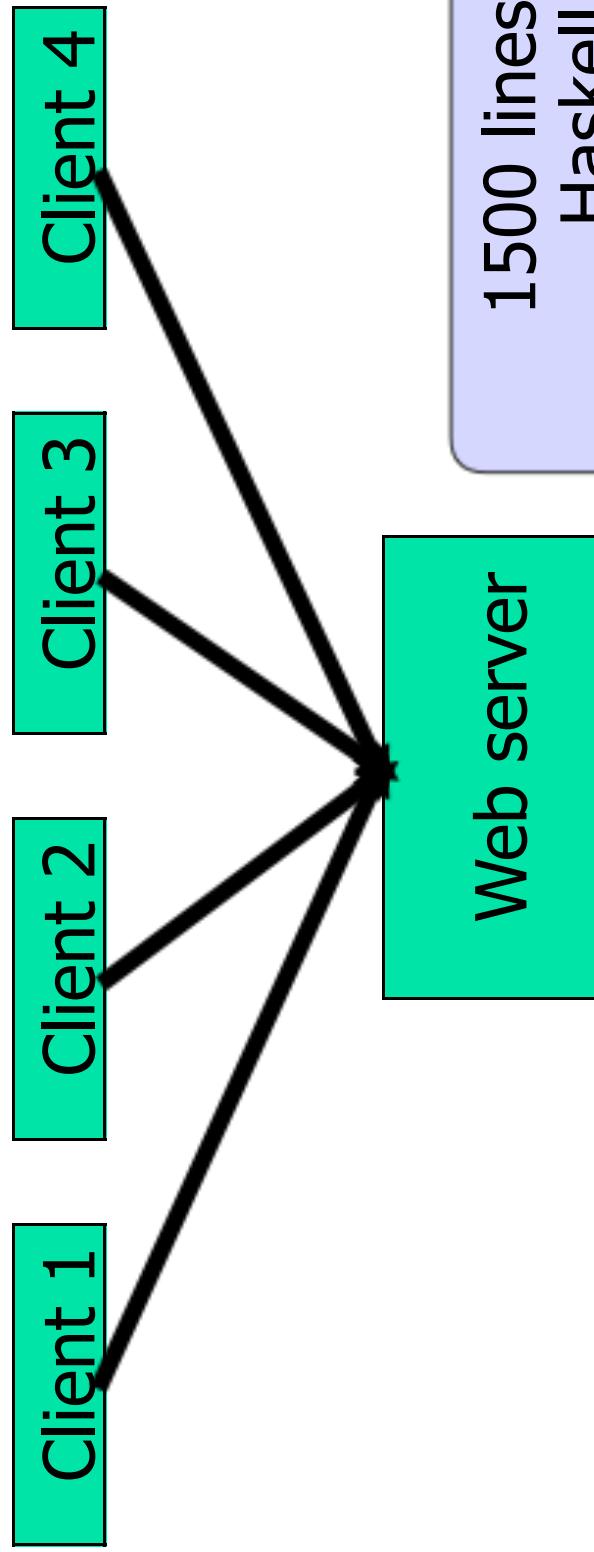
Tackling the awkward squad

- So lazy languages force us to take a different, more principled, approach to the Awkward Squad.
- I'm going to describe that approach for Haskell
- But I think there are lessons for strict, or even imperative, languages too.



A web server

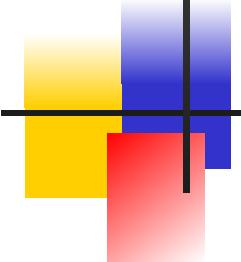
- I'll use a web server as my motivating example
- Lots of I/O, lots of concurrency, need for error recovery, need to call external libraries



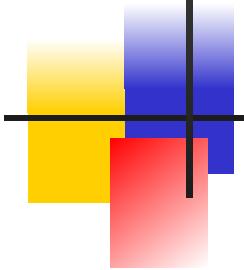
1500 lines of
Haskell
700
connections/sec

0

Monadic input and output



The problem

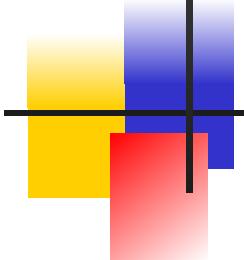


A functional program defines a pure function, with no side effects

Tension

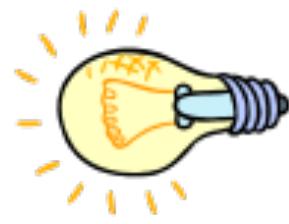
The whole point of running a program is to have some side effect

Functional I/O



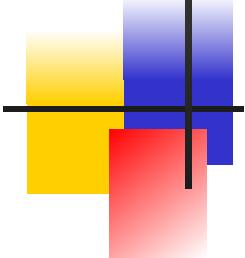
```
main :: [Response] -> [Request]
```

```
data Request = ReadFile Filename  
              | WriteFile FileName String  
              | ...
```



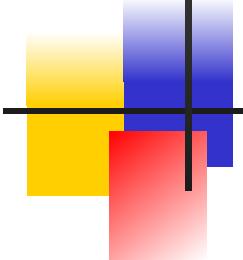
```
data Response = RequestFailed  
              | ReadOK String  
              | WriteOK  
              | ...
```

- “Wrapper program” interprets requests, and adds responses to input



Functional I/O is awkward

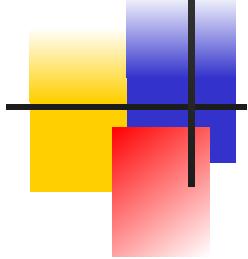
- Hard to extend (new I/O operations □ new constructors)
- No connection between Request and corresponding Response
- Easy to get “out of step” (□ deadlock)



Monadic I/O: the key idea

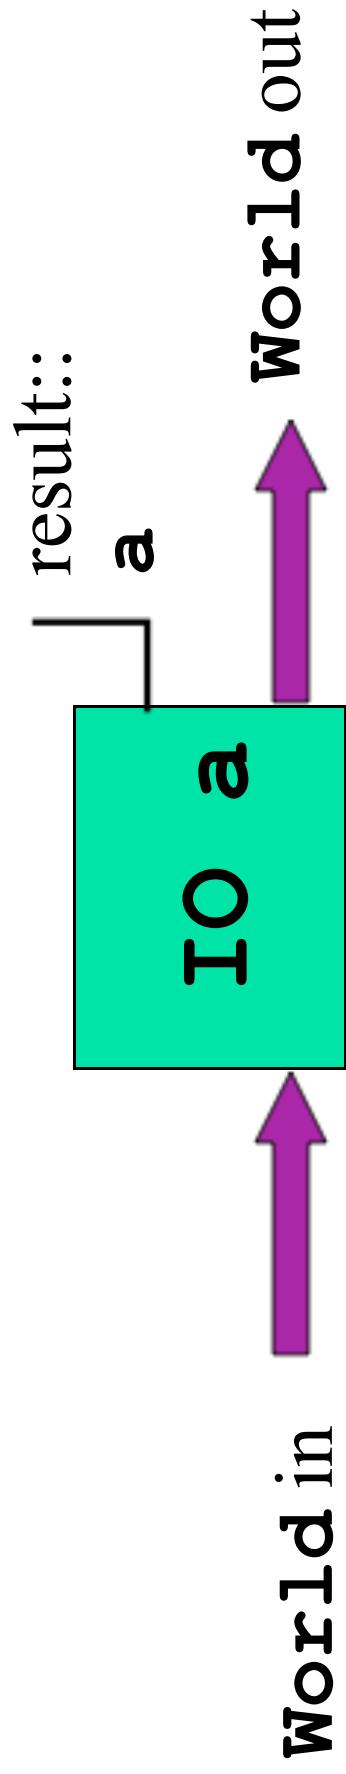
A value of type ($\text{IO } t$) is an “**action**” that, when performed, may do some input/output before delivering a result of type t .

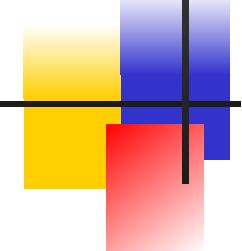
A helpful picture



A value of type (`IO t`) is an “**action**” that, when performed, may do some input/output before delivering a result of type `t`.

```
type IO a = World -> (a, World)
```





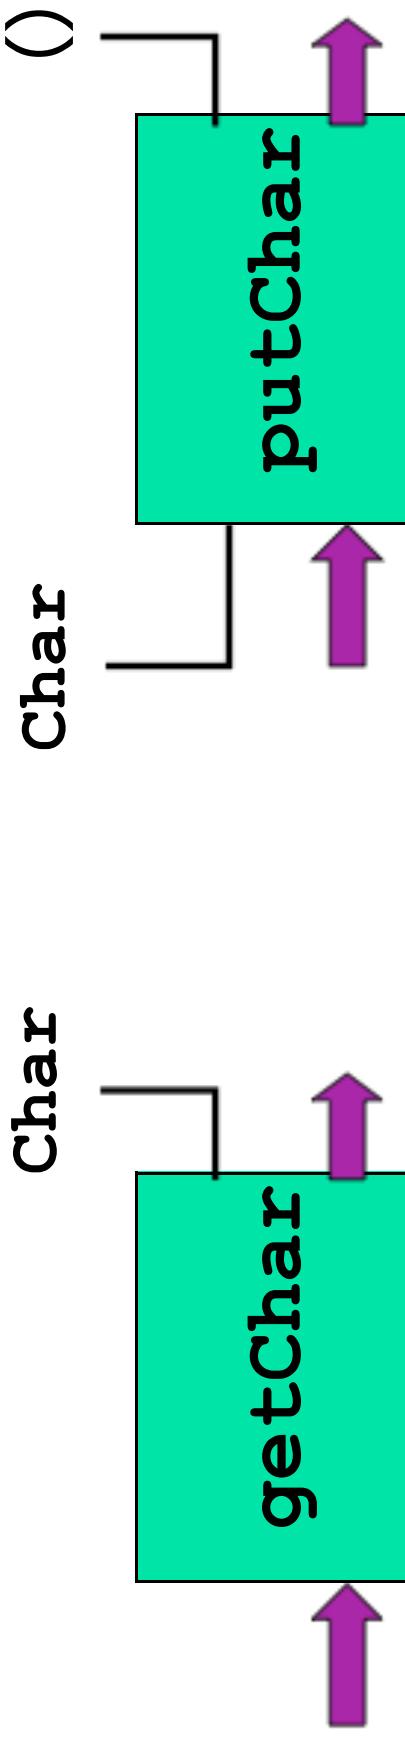
Actions are first class

A value of type (`IO t`) is an “**action**” that, when performed, may do some input/output before delivering a result of type `t`.

```
type IO a = World -> (a, World)
```

- “Actions” sometimes called “computations”
- An action is a **first class value**
- **Evaluating** an action has no effect; **performing** the action has an effect

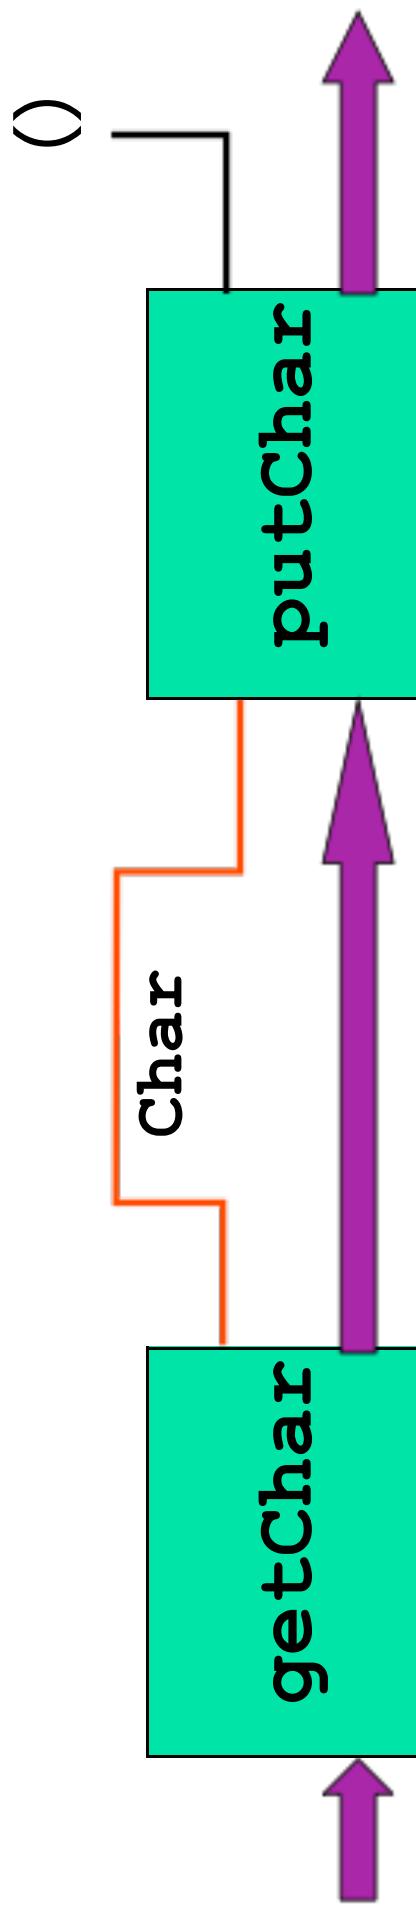
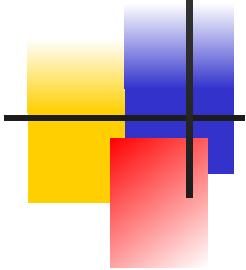
Simple I/O



```
getChar :: IO Char  
putchar :: Char -> IO ()  
  
main :: IO ()  
main = putchar 'x'
```

Main program is an action of type IO ()

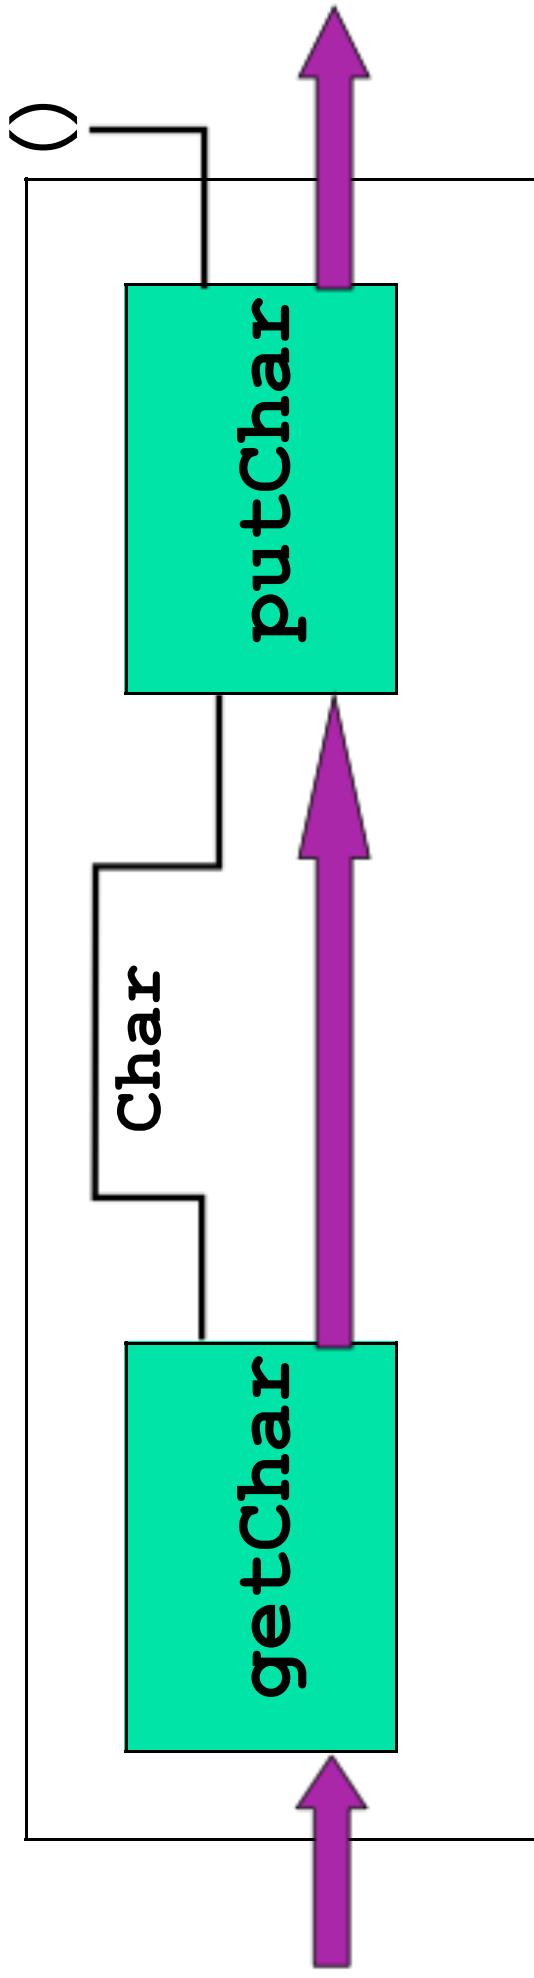
Connecting actions up



Goal: read a character and then write it back out

The ($>>=$) combinator

$(>>=) :: IO\ a \rightarrow (a \rightarrow IO\ b) \rightarrow IO\ b$



- We have connected two actions to make a new, bigger action.

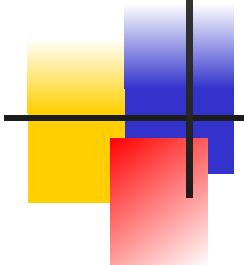
```
echo :: IO ()  
echo = getChar >>= putChar
```

Printing a character twice

```
echoDup :: IO ()  
echoDup = getChar >=> (\\c ->  
    putChar c <=> (\\()  
        ->  
        putChar c))
```

- The parentheses are optional

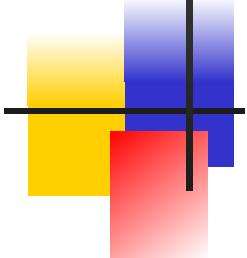
The ($\wedge\wedge$) combinator



```
echoDup :: IO ()  
echoDup = getChar <=> \c ->  
    putStrLn c  
    putStrLn c
```

```
(<>) :: IO a -> IO b  
m <> n = m && n  
a <> b = (x \n)
```

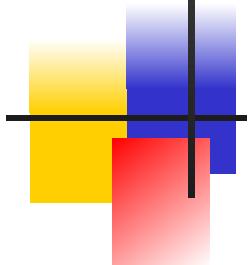
Getting two characters



```
getTwoChars :: IO (Char, Char)
getTwoChars = >>= \c1
->
getChar    >>= \c2 ->
                c2
```

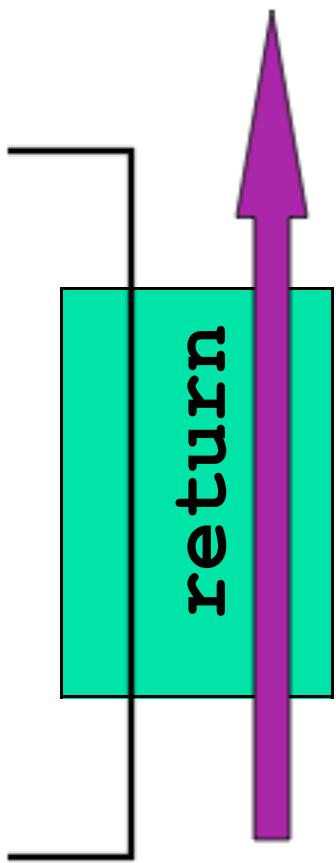
We want to just return (c1,c2)

The **return** combinator

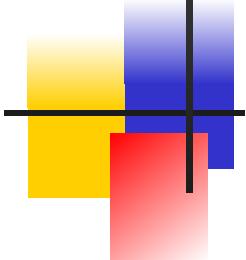


```
getTwoChars :: IO (Char, Char)
getTwoChars = getChar >>= \c1
              ->
              getChar >>= \c2 ->
return (c1, c2)
```

```
return :: a -> IO a
```



Notational convenience



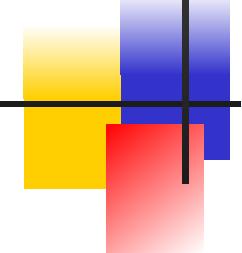
```
getTwoChars :: IO (Char,Char)
getTwoChars = getChar >>= \c1 ->
               getChar >>= \c2 ->
               return (c1,c2)
```

- By design, the layout looks imperative
 - c1 = getChar();
 - c2 = getChar();
 - return (c1,c2);

Notational convenience

```
getTwoChars :: IO (Char,Char)
getTwoChars = do { c1 <- getChar
                  ; c2 <- getChar
                  ; return (c1,c2) }
```

“do” notation adds only syntactic sugar



Desugaring do notation

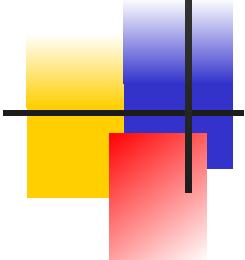
“do” notation adds only syntactic sugar

```
do { x<-e; s } = e >>= \x -> do { s }
do { e; s } = e >> do { s }
do { e } = e
```

Getting a line

```
getLine :: IO [Char]
getLine = do { c <- getChar;
              if c == '\n' then
                return []
              else do { cs <- getLine;
                        return (c:cs)
                      }
            }
```

Note the nested "do" expression



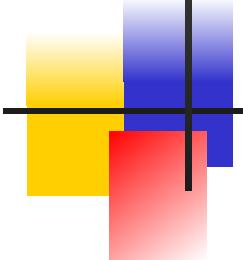
Control structures

Values of type (`IO#(T)`) are first class
So we can define our own "control structures"

```
forever :: IO#() -> IO#()
forever a = a >> forever a

repeatN :: Int -> IO#() -> IO#()
repeatN#(0) a = return()
repeatN#(n) a = a >> repeatN#(n-1) a
```

e.g. `repeatN#(10)(putChar('x'))`



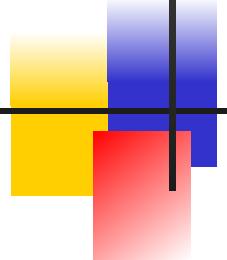
Loops

Values of type (`IO t`) are first class
So we can define our own "control structures"

```
for :: [a] -> (a -> IO b) -> IO ()  
for [] fa = return ()  
for (x:xs) fa = fa x >> for xs fa
```

e.g. `for [1..10] (\x -> putStrLn (show x))`

Loops



A list of IO actions

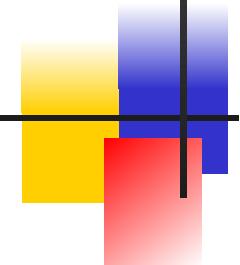
An IO action returning a list

```
sequence :: [IO a] -> IO [a]
sequence [] = return []
sequence (a:as) = do { r <- a;
                      rs <- sequence as;
                      return (r:rs) }
```

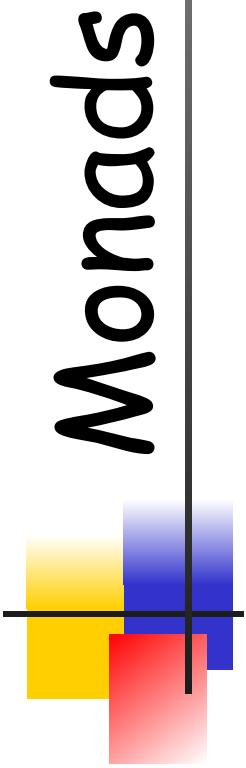


```
for :: [a] -> (a -> IO b) -> IO ()
for xs fa = sequence (map fa xs) >>
             return ()
```

First class actions



Slogan: first class actions
let us write application-specific control structures



Monads

A monad consists of:

A type constructor M
bind :: $M a \rightarrow (a \rightarrow M b) \rightarrow M b$
unit :: $a \rightarrow M a$

-
-
-

PLUS some laws

Laws for monads

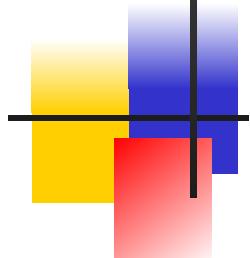
$$\text{unit } x \gg= f \quad = \quad f \times$$

$$m \gg= \text{unit} \quad = \quad m$$

$$m1 \gg= (\square x.m2) \gg= (\square y.m3))$$

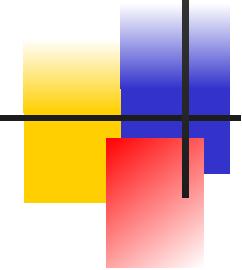
$$(m1 \gg= (\square x.m2)) \gg= (\square y.m3)$$

Third law in do-notation



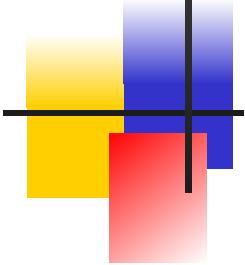
```
do { x <- do { y <- m2 ; }  
     m1 } ;  
     m3 } ;  
     m3 }
```

$$\begin{aligned}m1 >>= (\Box x.m2) >>= (\Box y.m3)) \\= \\(m1 >>= (\Box x.m2)) >>= (\Box y.m3)\end{aligned}$$



What does it all mean?

Functional semantics



`type IO a = World -> (a, World)`



- A program that loops forever has meaning □
A program that prints 'x' forever has meaning □
- What is the meaning of two Haskell programs running in parallel?
- Does not scale well to concurrency, non-determinism

Trace semantics

```
type IO a = Set Trace
type Trace = [Event]
type Event = ReadChar Char
| WriteChar Char
...
...
```

```
echo = { [ReadChar 'a'; WriteChar
'a'],
[ReadChar 'b'; WriteChar 'b'],
[ReadChar 'c'; WriteChar 'c'],
...}
```

- Can deal with non-determinism, and concurrency

P Q

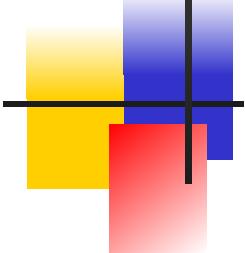
Operational semantics

Instead of saying what the meaning of a program **is**,
say how the program **behaves**

Equivalence of programs becomes **similarity of behaviour** instead of **identity of meaning**



Program state P can move to state Q, exchanging
event \square with the environment

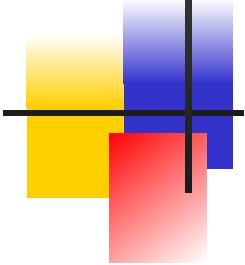


Program states

A **program state** represents the current internal state of the program.

Initially, it is just a term, $\{M\}$
Curly braces say
“here is a program state”

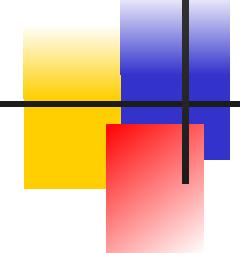
e.g. `{putChar 'x' >> putStrLn "y"}`



Events

Events describe how the program interacts with the external world: i.e. what I/O it performs

- $P \square Q$ P can move to Q, writing c to stdout
- $P \square_Q^c$ P can move to Q, reading c to stdin
- $P \square_{?c}$



Our first two rules

```
{putChar ch}      ! ch → {return ()}  
{getChar}        ? ch → {return ch}
```

But: what happens to this?

{getChar >>= \c -> putChar c} □ ???

Want to say "look at the action in the leftmost position"

Evaluation contexts

$E ::= [.] \quad | \quad E >>=$

M

An **evaluation context** E is a term with a “hole” in it.

For example:

$E1 = [.] >>= M$

$E2 = ([.] >>= M1) >>= M2$

$E[M]$ is the evaluation context E with the hole filled in by M . So

$E1[getChar] = getChar >>= M$
 $E2[getChar] = (getChar >>= M1) >>= M2$

Revised rules for put/get

$E ::= [.] \quad | \quad E >>=$

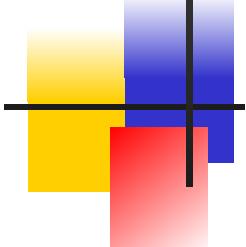
M

$$\begin{array}{ccc} \{E[\text{putChar } ch]\} & \xrightarrow{!_{ch}} & \{E[\text{return } ()]\} \\ \{E[\text{getChar}]\} & \xrightarrow{?_{ch}} & \{E[\text{return } ch]\} \end{array}$$

{getChar} $>>= \backslash c -> \text{putChar } c\}$

? $\{ \text{return } ch \} \xrightarrow{ch} \backslash c -> \text{putChar } c\}$

The return/bind rule

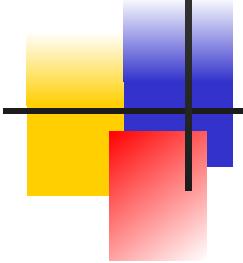


```
{E[return N>>= M]} → {E[M N]}
```

```
{getChar >>= \c-> putChar c}  
? █ {return ch >>= \c-> putChar c}  
ch █ {(\c-> putchar c) ch}
```

Silent
transition

Now we need to do some "ordinary evaluation"



The evaluation rule

$\mathcal{E}[M]$ is the denotational semantics for M

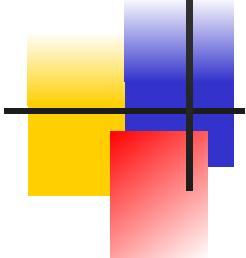
V is the value of M

M wasn't already evaluated

If the things above the line are true, then we can deduce the thing below the line

$$\frac{\mathcal{E}[M] = V \quad M \not\equiv V}{\{E[M]\} \rightarrow \{E[V]\}}$$

The evaluation rule



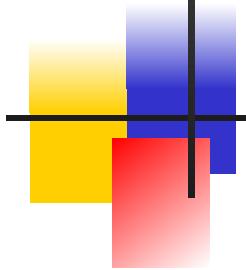
$$\frac{\mathcal{E}[M] = V \quad M \not\equiv V}{\{\mathcal{E}[M]\} \rightarrow \{\mathcal{E}[V]\}}$$

$\mathcal{E}[M]$ is the denotational semantics for M

② $\{(\backslash c \rightarrow \text{putChar } c) \text{ } ch\}$
③ $\{\text{putChar } ch\}$
 $!c$ ③ $\{\text{return } ()\}$
 h

Treat primitive IO actions as “constructors”; so $(\text{putChar } ch)$ is a value

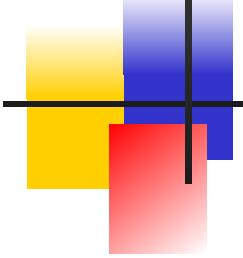
Mutable variables in C



```
int x =  
3;  
x = x+1;
```

x is a
location initialised
with 3

read x, add 1,
store into x



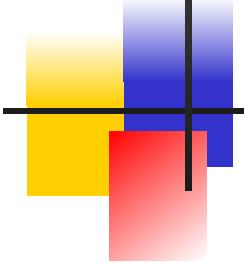
Mutable variables in Haskell

```
do { x <- newIORef 3;  
    v <- readIORef x;  
    writeIORef x (v+1) }
```

x is a location,
initialised with 3

read x, add 1,
store into x

```
newIORef   :: a -> IO (IORef a)  
readIORef  :: IORef a -> IO a  
writeIORef :: IORef a -> a -> IO ()
```



Semantics for variables

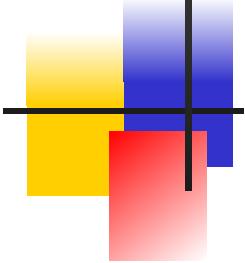
Step 1: elaborate the program state

$P, Q, R ::=$	$\{M\}$	The main program
	$\langle M \rangle_r$	An IORef named r , holding M
	$P \mid Q$	Parallel composition
	$\nu x.P$	Restriction

e.g. $\Box r.s. (\{M\} \mid \langle 3 \rangle_r \mid \langle 89 \rangle_s)$

The main program

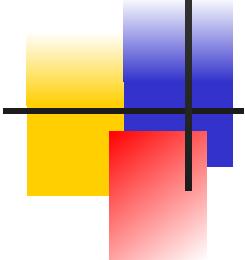
An IORef called r ,
holding 3



Semantics for variables

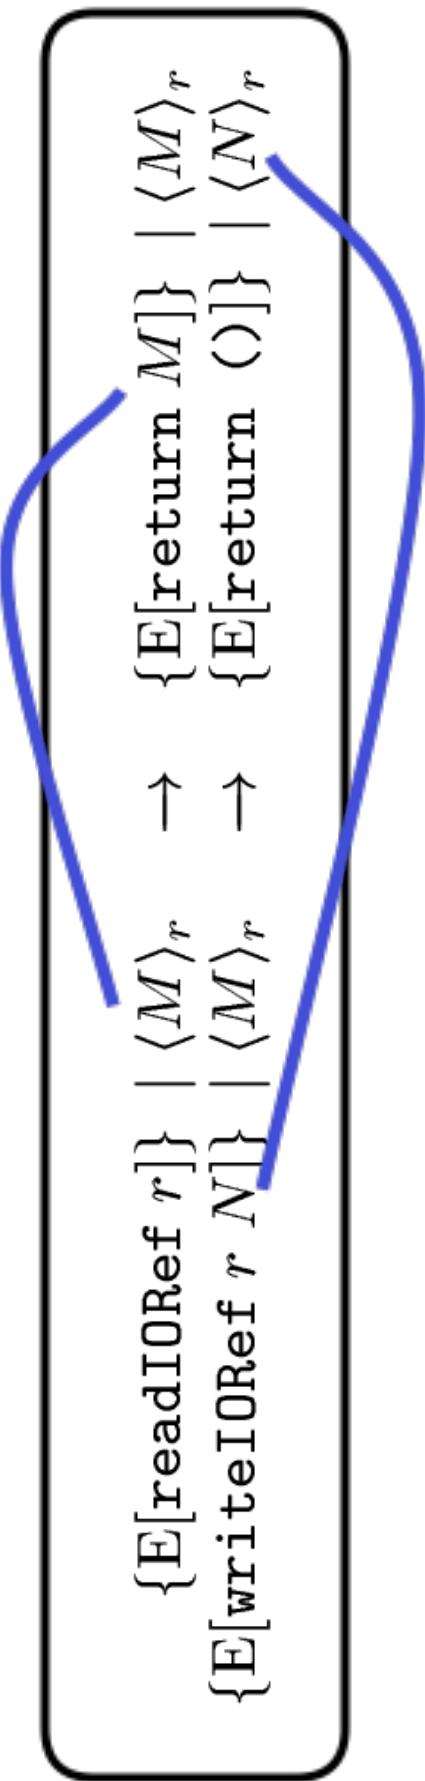
Step 2: add rules for read, write IORefs

$$\begin{array}{l|l} \{E[\text{return } M]\} & | \\ \{E[\text{return } N]\} & | \\ \{E[\text{readIORef } r]\} & | \langle M \rangle_r \\ \{E[\text{writeIORef } r\ N]\} & | \langle M \rangle_r \end{array} \rightarrow \begin{array}{l|l} \{E[\text{return } M]\} & | \langle M \rangle_r \\ \{E[\text{return } O]\} & | \langle N \rangle_r \end{array}$$

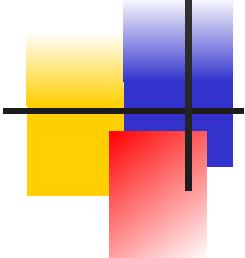


Semantics for variables

Step 2: add rules for read, write IORefs



But what if the main program is not “next to” the relevant variable?



"Structural rules"

Can stir the mixture

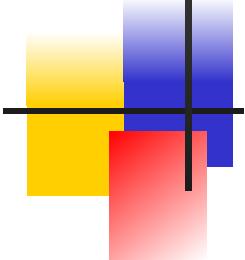
Step 3: add rules to bring "reagents" together

$$P \mid (Q \mid R) \equiv Q \mid P \quad (P \mid Q) \mid R$$

$$\frac{P \equiv P' \quad P' \xrightarrow{\alpha} Q' \quad Q' \equiv Q}{P \xrightarrow{\alpha} Q} \quad (EQUIV)$$

$$\frac{P \xrightarrow{\alpha} Q}{P \mid R \xrightarrow{\alpha} Q \mid R} \quad (PAR)$$

Can look under "|"

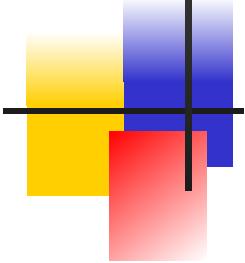


Restriction

Step 4: deal with **fresh IORef names**

{E [newIORef M] } {E [return ?] } | <M>?

What can we use
as the IORef
name???



Restriction

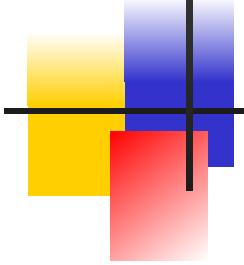
Step 4: deal with **fresh IORef names**

r is not used
already

$$\frac{r \notin fn(E, M)}{\{E[\text{newIORef } M]\} \rightarrow \nu r.(\{E[\text{return } r]\} \mid \langle M \rangle_r)}$$

Restriction

More "Structural rules"



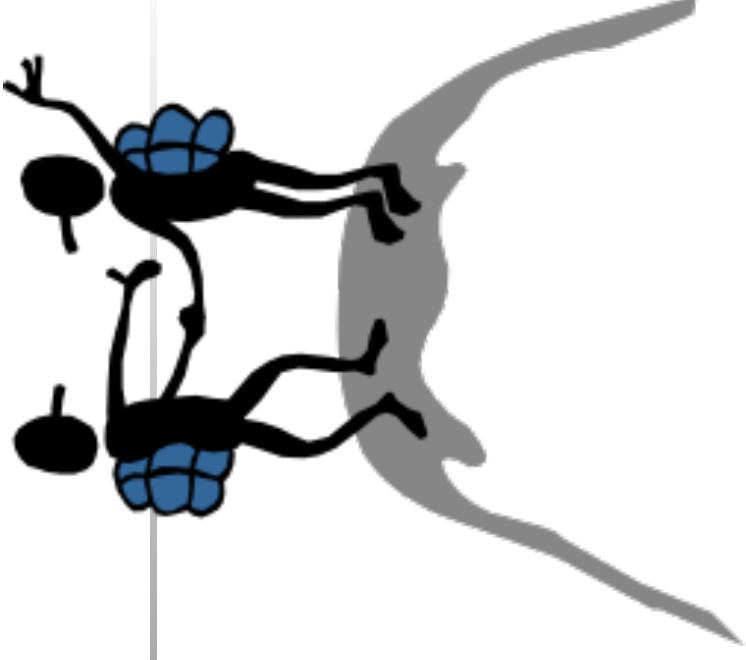
Step 5: **structural rules for restriction**

Can float
“ \Box ”
outwards

$$\begin{array}{ccl} \nu x.\nu y.P & \equiv & \nu y.\nu x.P \\ (\nu x.P) \mid Q & \equiv & \nu x.(P \mid Q), \quad x \notin fn(Q) \\ \nu x.P & \equiv & \nu y.P[y/x], \quad y \notin fn(P) \end{array}$$

$$\frac{P \xrightarrow{\alpha} Q}{\nu x.P \xrightarrow{\alpha} \nu x.Q} (NU)$$

Can look
under “ \Box ”



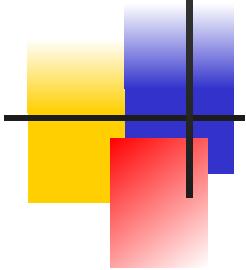
Phew!

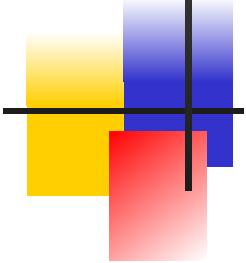
Quite a lot of technical machinery! But:

It's standard, widely-used, machinery (esp in process calculi), so it's worth getting used to

- It scales to handle non-determinism, concurrency
- The machinery (esp structural rules) doesn't get in the way most of the time

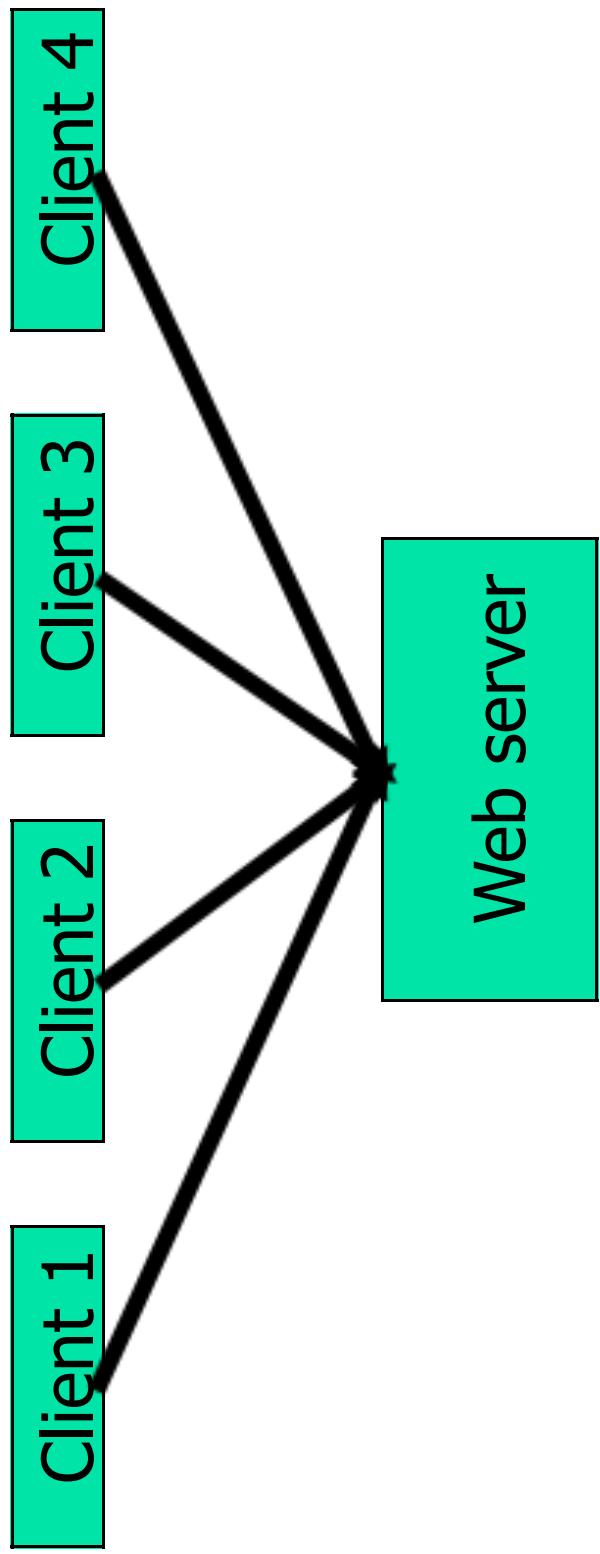
Concurrency

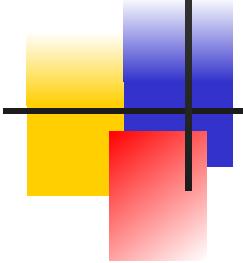




The need for concurrency

- Want one thread ("virtual server") per client
- Threads largely independent, but share some common resources (e.g. file system)





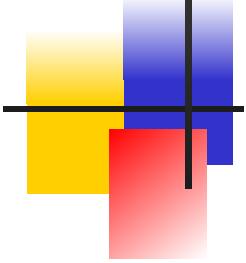
Concurrency vs parallelism

Parallel functional programming:

- Aim = performance through multiple processors (e.g. e_1+e_2 in parallel)
- No semantic changes; deterministic results

Concurrent functional programming

- Aim = concurrent, **I/O-performing** threads
- Makes perfect sense on a uniprocessor
- Non-deterministic interleaving of I/O is inevitable

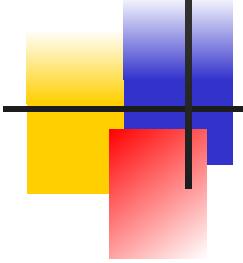


Concurrent web service

```
acceptConnections :: Config -> Socket -> IO ()  
acceptConnections config socket  
= forever (do { conn <- accept socket ;  
    forkIO (serviceConn config conn) })
```

forkIO :: IO a -> IO ThreadId

- forkIO spawns an *independent*, I/O performing, thread
- No parameters passed; free variables do fine



Communication and sharing

- What if two threads want to communicate?
 - Or share data?
- Example: keep a global count of how many client threads are running
 - Increment count when spawning
 - Decrement count when dying



Communication and sharing

```
data MVar a          :: IO (MVar a)
newEmptyMVar        :: MVar a -> a -> IO ()
putMVar             :: MVar a -> IO a
takeMVar            :: MVar a -> IO a
```

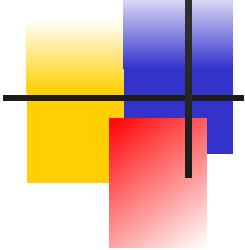
- A value of type `(MVar t)` is a **location** that is either
 - **empty**, or
 - **holds a value** of type `t`

27

particular
Maven
dependency
Leave

Maven dependency
Leave

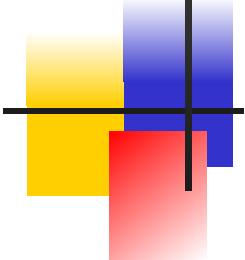
Communication and sharing



Using MVars

```
acceptConnections :: Config -> Socket -> IO ()  
acceptConnections config socket  
= do { count <- newEmptyMVar ;  
      putMVar count 0 ;  
      forever (do { conn <- accept socket ;  
                    forkIO (do { inc count ;  
                               serviceConn config conn ;  
                               dec count} ) }  
  
inc,dec :: MVar Int -> IO ()  
inc count = do { v <- takeMVar count ; putMVar count (v+1) }  
dec count = do { v <- takeMVar count ; putMVar count (v-1) }
```

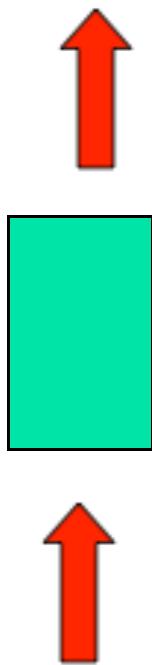
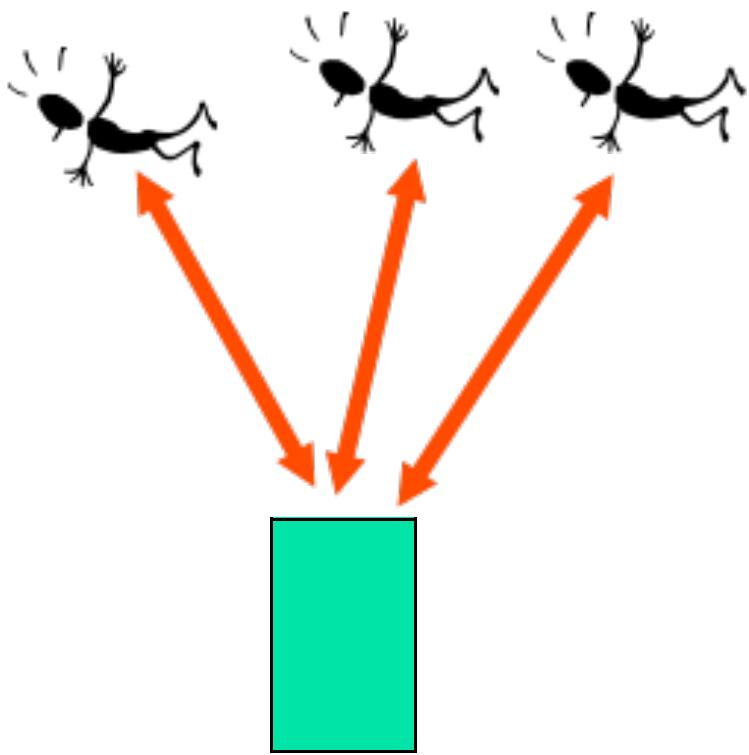
MVar is empty at this point, hence no race hazard

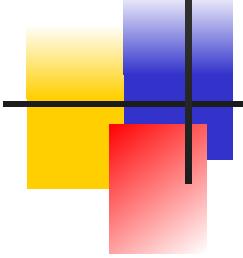


MVars as channels

An MVar directly implements:

- a shared data structure
- a one-place channel
-





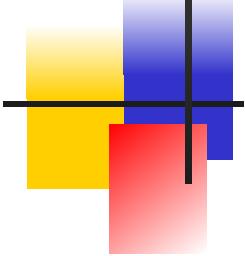
Semantics

Fortunately, most of the infrastructure is there already!

Step 1: elaborate the program state

$P, Q, R ::= \dots$		
	$\{M\}_t$	A thread called t
	$\langle M \rangle_m$	An MVar called m containing M
	$\langle \rangle_m$	An empty MVar called m

e.g. `{putChar 'c' } t1 | {putChar 'd' } t2`



Semantics

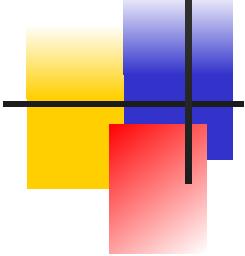
Step 2: a rule for `forkIO`

$$\frac{\{E[\text{forkIO } M]\}_t \rightarrow \nu u.(\{E[\text{return } u]\}_t \mid \{M\}_u)}{u \notin fn(M, E)} (FORK)$$

Restrict the new thread name

Return the thread name to the caller

The new thread

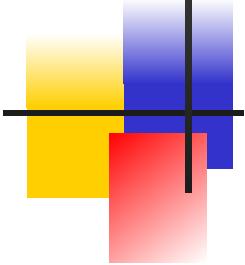


Semantics

Step 3: **a rules for new, take, put**

$$\begin{array}{lcl} \{E[takeMVar\ m]\}_t \mid \langle M \rangle_m & \rightarrow & \{E[return\ M]\}_t \mid \langle \rangle^m \\ \{E[putMVar\ M]\}_t \mid \langle \rangle_m & \rightarrow & \{E[return\ ()]\}_t \mid \langle M \rangle_m \end{array}$$

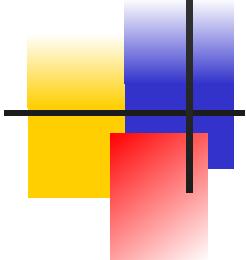
- Same as `readIORef`, `writeIORef`, except that `MVar` is filled/emptied
 - Blocking is implicit
 - Non-determinism is implicit



Building abstractions

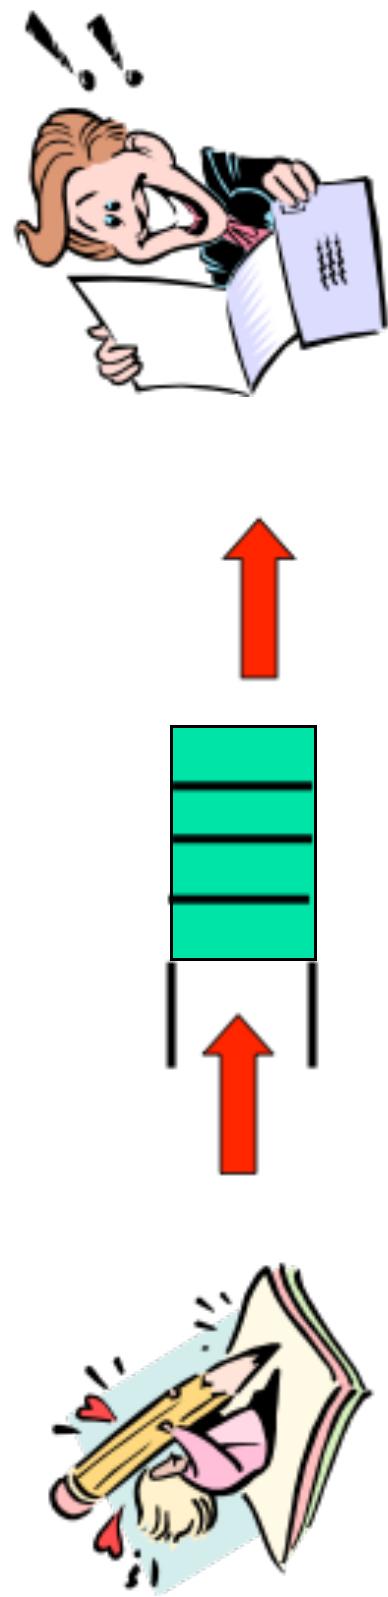
- MVars are primitive
- Want to build abstractions on top of them
- Example: a buffered channel

A buffered channel



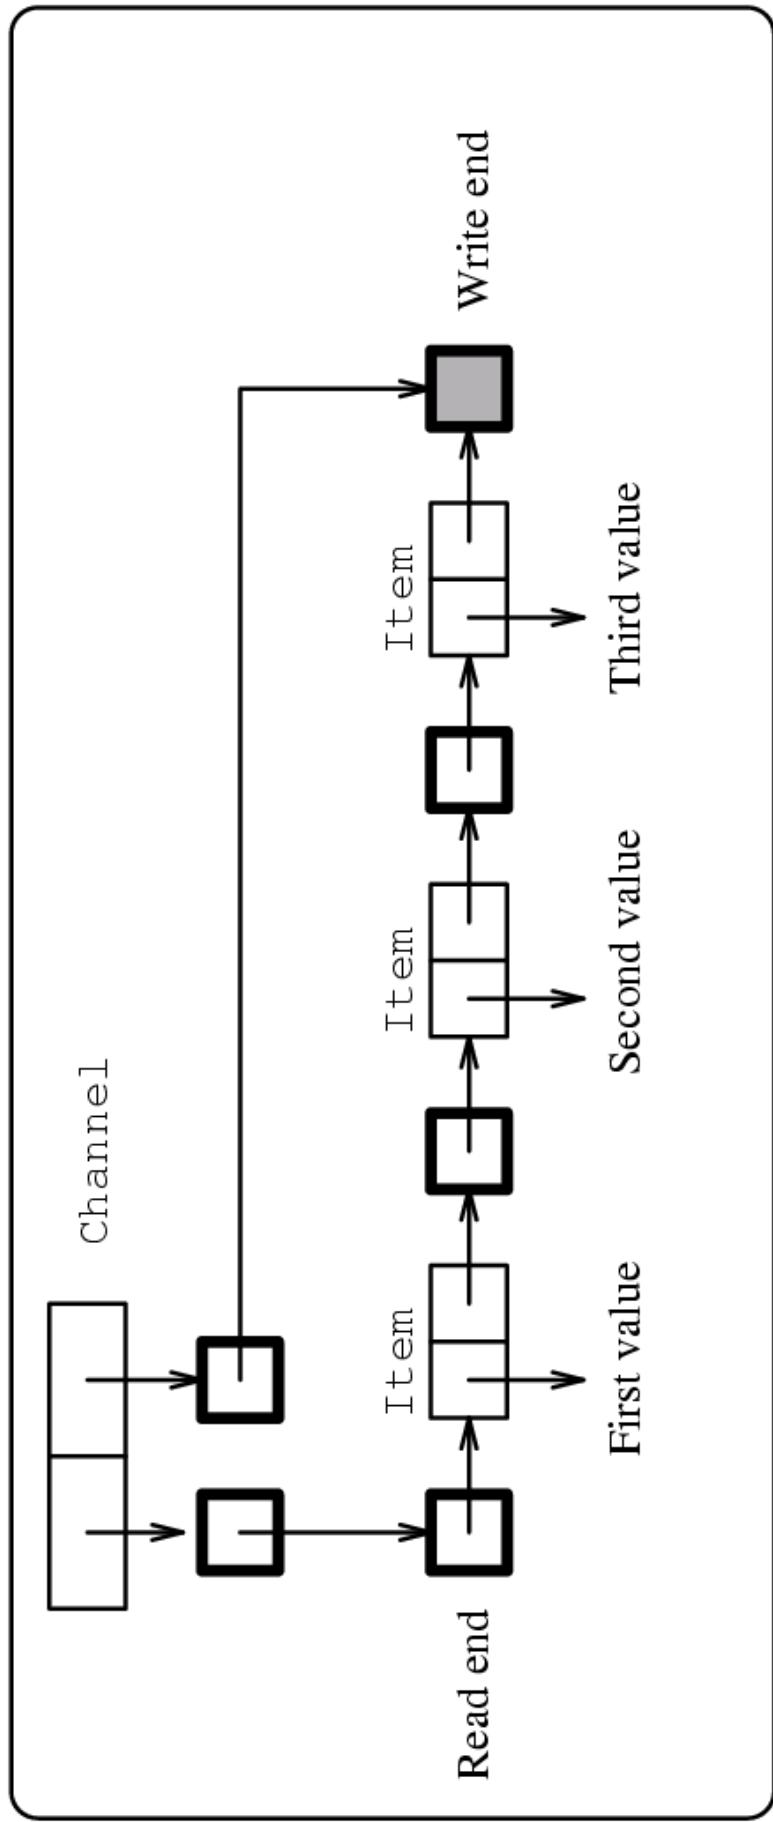
```
type Chan a :: IO (Chan a)
newChan :: Chan a -> a -> IO a
putChan :: Chan a -> a -> IO a
()
```

```
getChan :: Chan a -> IO a
```



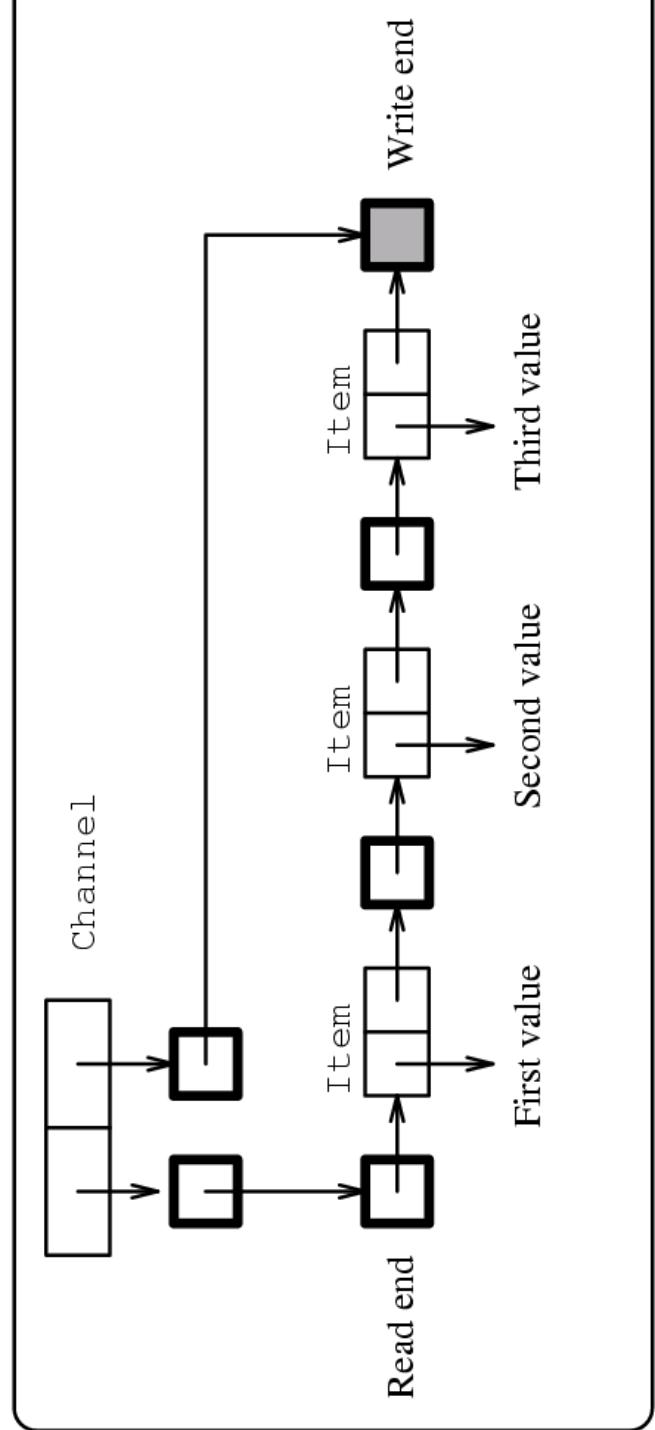
A buffered channel

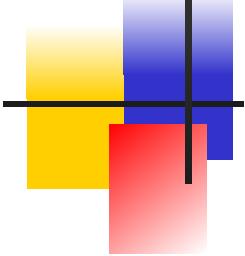
```
type Chan a    = (MVar (Stream a) ,MVar (Stream a))
type Stream a = MVar (Item a)
data Item a   = MkItem a (Stream a)
```



A buffered channel

```
putChan :: Chan a -> a -> IO ()  
putChan (read,write) val  
= do { new_hole <- newEmptyVar ;  
      old_hole <- takeMVar write ;  
      putMVar write new_hole  
      putMVar old_hole (MkItem val new_hole) }
```

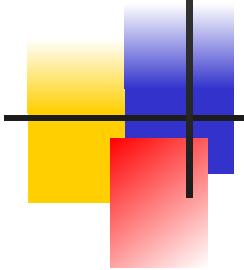


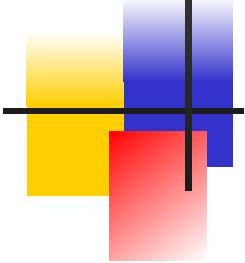


Summary

- `forkIO + MVars` are very simple.
- MVars are primitive, but surprisingly often Just The Right Thing
- Other excellent references
 - Concurrent programming in ML (Reppy, CUP)
 - Concurrent programming in Erlang (Armstrong, Prentice Hall, 2nd edition)

Exceptions



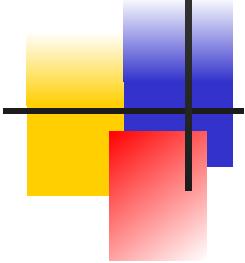


Why do we need exceptions?

Robust programs deal gracefully with "unexpected" conditions. E.g.

- o Disk write fails because disk is full
- o Client goes away, so server should time out and log an error
- o Client requests seldom-used service; bug in server code gives pattern-match failure or divide by zero

Server should not crash if these things happen!

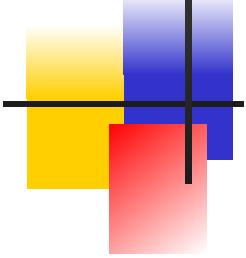


Approach 1: virtue

"A robust program never goes wrong"
(e.g. test for disk full before writing)

BUT:

- Can't test for all errors (e.g. timeouts)
 - Some problems are, by definition, errors in the server code
- Need a way to recover from ANY error**



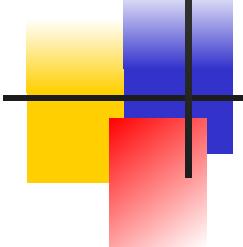
Approach 2: exceptions

Provide a way to say "execute this code, but if anything (at all) goes wrong, abandon it and do this instead".

I call this

"Exceptions for disaster recovery"

- Exception handler typically covers a large chunk of code
- Recovery action typically aborts a whole chunk of work



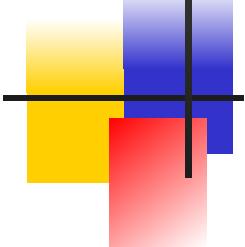
Aside: bad uses of exceptions

Exceptions are often (mis-) used in a different way:

"Exceptions for extra return values"

e.g. Look up something in a table, raise "NotFound" if it's not there.

- Exception handler often encloses a single call
 - Recovery action typically does not abort anything



Exceptions in Haskell 98

Haskell 98 supports exceptions in I/O actions:

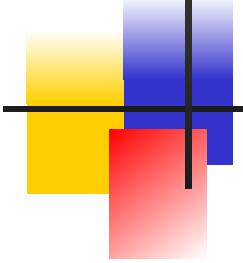
```
catch    :: IO a -> (IOError -> IO a) -> IO a  
userError :: String -> IOError  
ioError   :: IOError -> IO a
```

```
catch    (do { h <- openFile  
             ``foo'';  
           processFile h })  
(\e -> putStrLn "Oh dear")
```

Protected
code

Exception
handler

Dynamic scope: exceptions raised in
processFile are caught

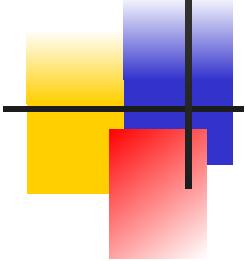


Semantics

Step 1: add a new *evaluation context*

$E ::= [.] \quad | \quad E \gg= M \quad | \quad \text{catch } E$

Says: "evaluate
inside the first
argument of
catch"



Semantics

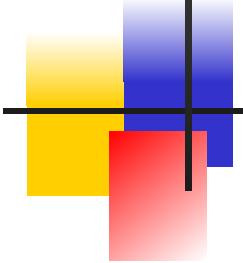
Step 2: add propagation rule for `ioError`

$$\{E[\text{ioError } e >>= M]\}_t \rightarrow \{E[\text{ioError } e]\}_t$$

An exception before
the (`>>=`)...

...discards the part
after the (`>>=`)

Standard stack-unwinding implementation is possible



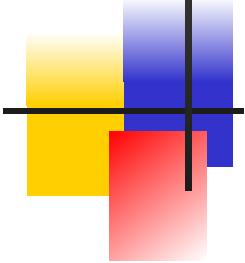
Semantics

What to do if an exception *is* raised

Step 3: add rules for catch

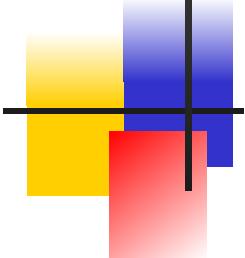
$$\begin{array}{ccc} \{E[\text{catch (ioError } e) M]\}_t & \rightarrow & \{E[M\ e]\}_t \\ \{E[\text{catch (return } N) M]\}_t & \rightarrow & \{E[\text{return } N]\}_t \end{array}$$

What to do if an exception *is not* raised



Synchronous vs asynchronous

- A **synchronous exception** is raised as a direct, causal result of executing a particular piece of code
 - Divide by zero
 - Disk full
- An **asynchronous exception** comes from "outside" and can arrive at any moment
 - Timeout
 - Stack overflow



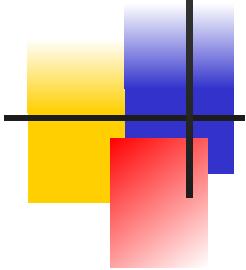
Haskell 98 isn't enough

Haskell 98 deals only with **synchronous exceptions** in the **IO monad**

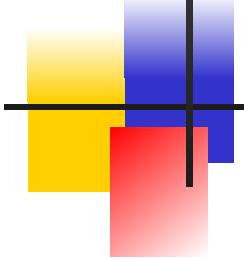
Two big shortcomings

- Does not handle things that go wrong in **purely-functional code**
- Does not deal with **asynchronous exceptions**

Exceptions in
pure code



Embed exceptions in values



Idea: embed exceptions in values

```
throw :: Exception -> a  
result type unchanged
```

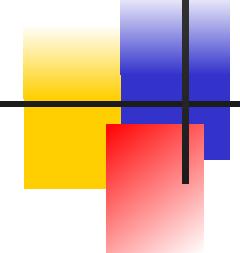
```
divide :: Int -> Int -> Int  
divide x y = if y==0 then throw  
DivZero  
else x/y
```



A value is

- either an “ordinary” value
- or an “exception” value, carrying an exception (Just like NaNs in IEEE floating point.)

In a lazy language an exception value might hide inside an un-evaluated data structure, but that's OK.



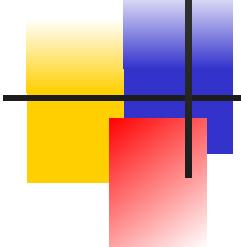
catching exceptions

New primitive for catching exceptions: BAD BAD!

```
getException :: a -> ExVal a  
  
data ExVal a = OK a  
             | Bad Exception
```

Example

```
f x = case getException (goop x) of  
        OK result -> result  
        Bad exn -> recovery_goop x
```



A well-known problem

What exception is raised by "+?"

(**throw ex1**) + (**throw**

ex2)

Usual **answer**: fix evaluation order

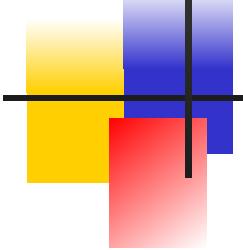
BAD ENOUGH for call-by-value languages

- loss of code-motion transformations
- need for effect analyses

TOTAL CATASTROPHE for Haskell

- evaluation order is deliberately unspecified
- key optimisations depend on changing evaluation order



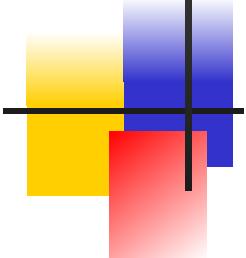


A cunning idea

Return both exceptions!

- A value is
 - either a "normal value"
 - or an "exceptional value"
- containing a **set** of exceptions

- Operationally, an exceptional value is
 - represented by a single representative
 - implemented by the usual stack-unwinding stuff
- c.f. infinite lists:
semantically infinite, operationally finite



Semantics without exceptions

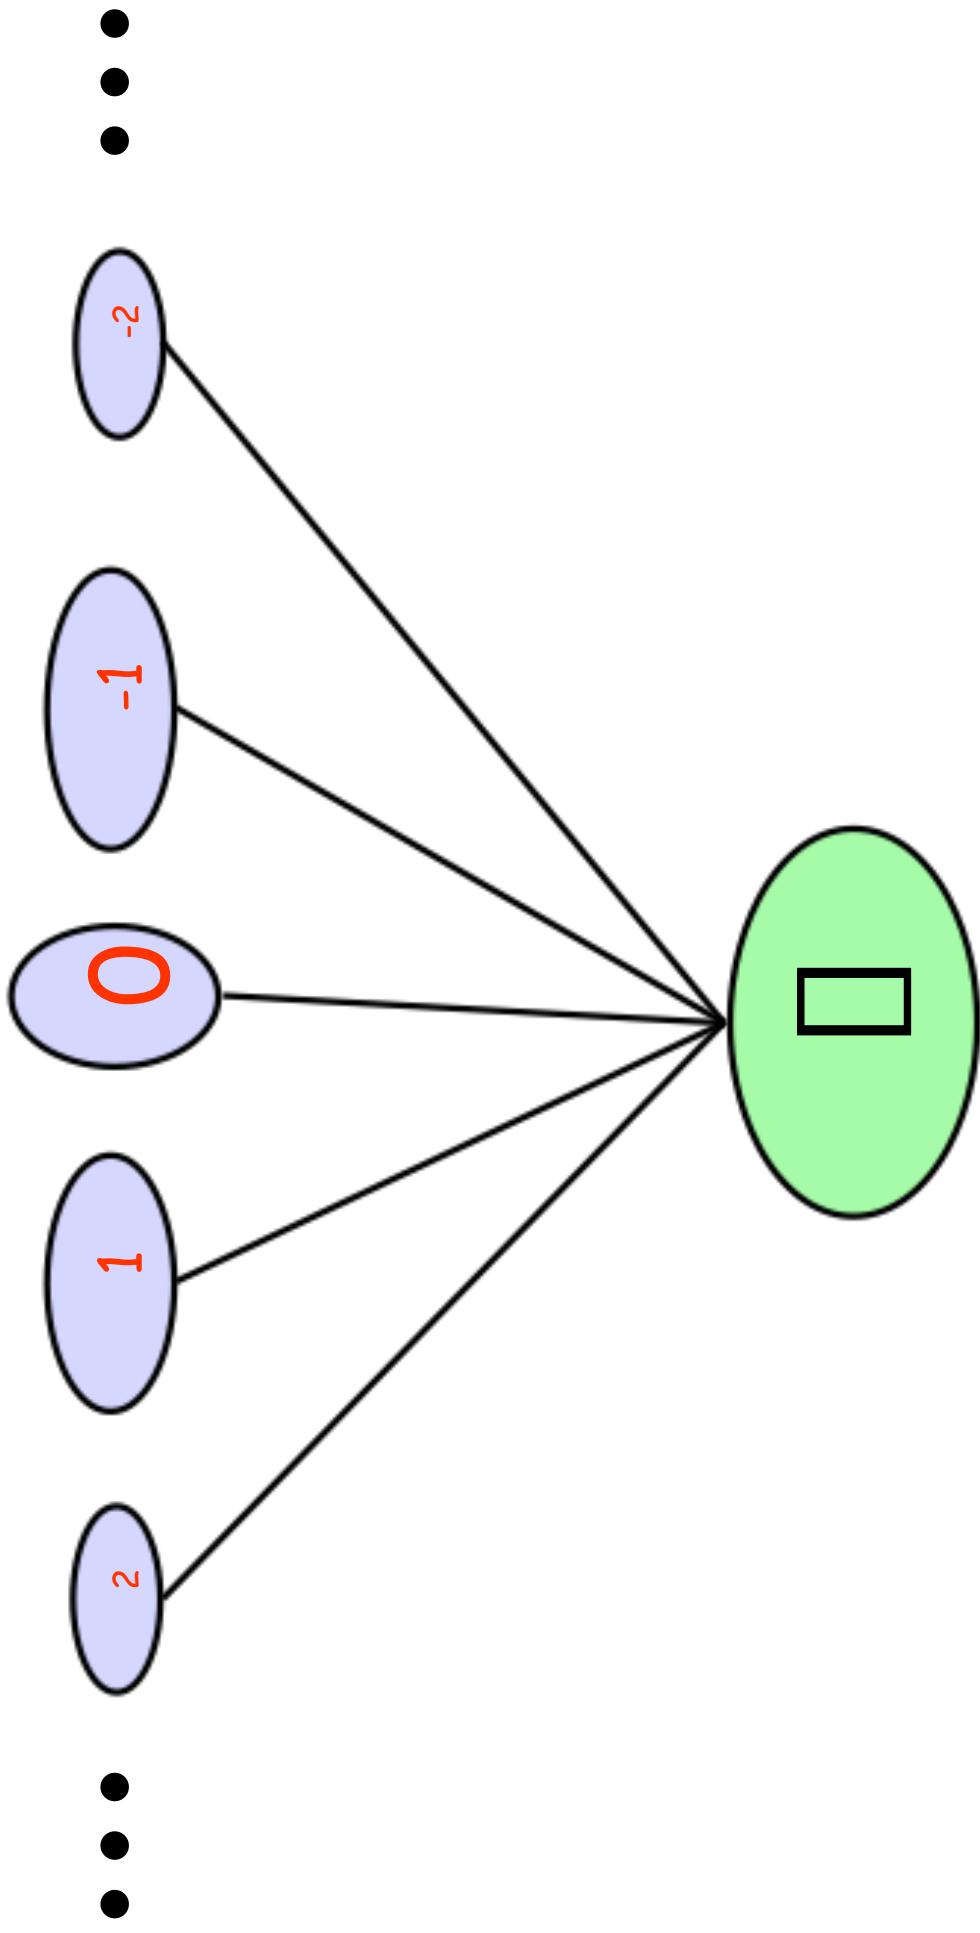
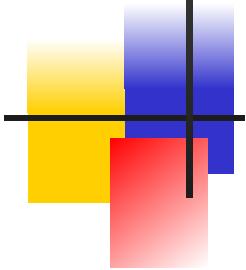
Denotations of Haskell types, $\llbracket T \rrbracket$

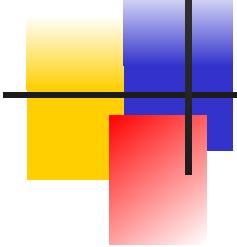
$$\begin{aligned}\llbracket \text{Int} \rrbracket &= \square M Z \\ \llbracket t_1 \Box t_2 \rrbracket &= \square M(\llbracket t_1 \rrbracket \Box \llbracket t_2 \rrbracket) \\ \llbracket (t_1, t_2) \rrbracket &= \square M(\llbracket t_1 \rrbracket \times \llbracket t_2 \rrbracket)\end{aligned}$$

$$M \dagger = + \Box \{\Box\}$$

e.g. $\llbracket \text{Int} \rightarrow \text{Int} \rrbracket = M(M \text{ Int} \rightarrow M \text{ Int})$

[[Int]] : Ordinary Haskell





Semantics with exceptions

Denotations of Haskell types, $\llbracket T \rrbracket$

Same as before

$$\begin{aligned}\llbracket \text{Int} \rrbracket &\quad \square M Z \\ \llbracket t_1 \Box t_2 \rrbracket &\quad \square M (\llbracket t_1 \rrbracket \Box \llbracket t_2 \rrbracket) \\ \llbracket (t_1, t_2) \rrbracket &\quad \square M (\llbracket t_1 \rrbracket \times \llbracket t_2 \rrbracket)\end{aligned}$$

Bottom

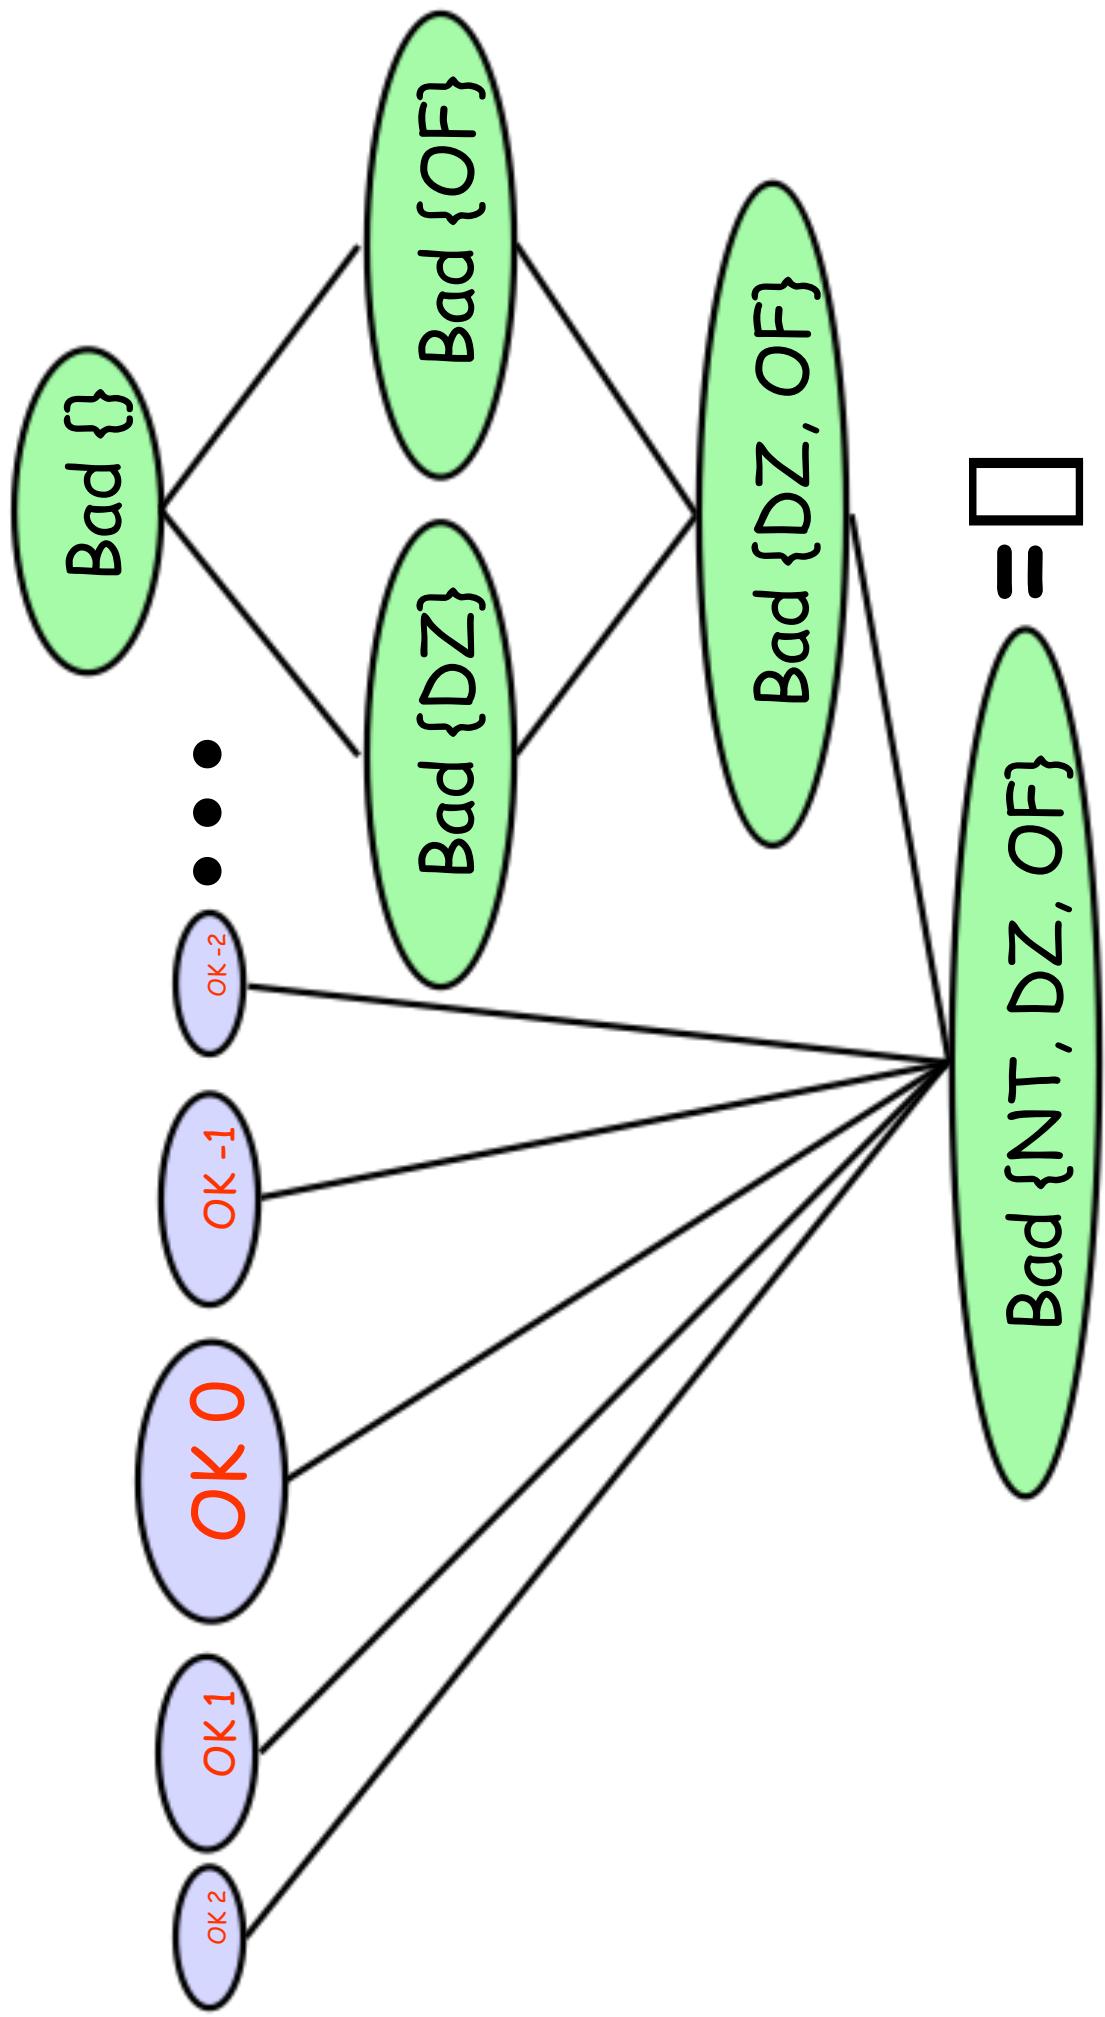
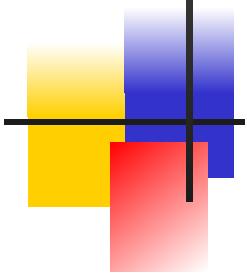
$$\begin{aligned}M + &= \{ \text{Ok} \times | \times \Box + \} \Box \\ \{ \text{Bad } s \mid s \Box E \} &\Box \\ \{ \text{Bad } (E \Box \{\text{NonTermination}\}) \} &\end{aligned}$$

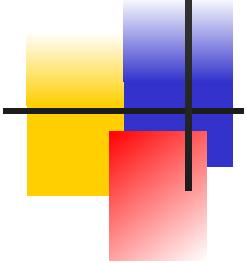
E $\square \{\text{Overflow}, \text{DivZero}, \dots\}$ (fixed set)

Z \square the integers

e.g. $\llbracket \text{Int} \rightarrow \text{Int} \rrbracket = M(M \text{ Int} \rightarrow M \text{ Int})$

[[Int]]: Exceptional Haskell



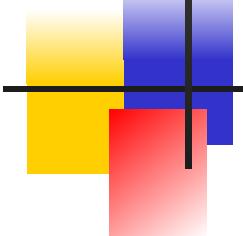


Semantics

$\llbracket e_1 + e_2 \rrbracket = \text{OK } (m + n)$ if $\text{OK } m = \llbracket e_1 \rrbracket$ \square
 $\text{OK } n = \llbracket e_2 \rrbracket$ \square
 $= \text{Bad}(\text{S}(\llbracket e_1 \rrbracket \square) \square \text{S}(\llbracket e_2 \rrbracket \square))$ otherwise

$$\begin{aligned}\text{S(Bad } s) &= S \\ \text{S(OK } n) &= \{\}\end{aligned}$$

Payoff: $\llbracket e_1 + e_2 \rrbracket = \llbracket e_2 + e_1 \rrbracket$



Whoa! What about `getException`?

Problem: which exception does `getException` choose from the set of possibilities?

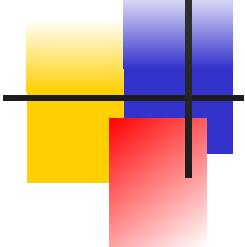


```
getException :: a -> ExVal a  
data ExVal a = OK a | BadException  
?????
```

Solution 1: **choose any.** But that makes `getException` non-deterministic. And that loses even \Box -reduction!

```
let x = getException e in x==x = True  
(getException e) == (getException e) □ True
```

Verdict: Cure worse than disease.

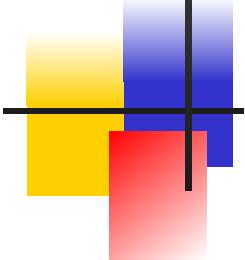


Using the IO monad

Solution 2: put getException in the IO monad:

```
evaluate :: a -> IO a
```

- `evaluate` evaluates its argument;
- if it is an ordinary value, it returns it
- if it is an exceptional value, it chooses one of the set of exceptions and raises it as an IO monad exception



Using the IO monad

Key idea:

The choice of which exception to raise is made in the IO monad, so it can be non-deterministic (like so much else in the IO monad)

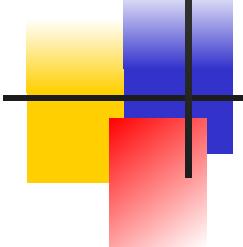
evaluate :: a -> IO a

Using evaluate

```
main = do {    i <- getInput;  
                
    catch (do {      r <- evaluate (goop i);  
                      do_good_stuff r  
                  }  
          (\ ex -> recover_from ex)  
    )  
}
```

You have to be in the IO monad to use evaluate

You do **not** have to be in the IO monad to use throw



Semantics

Add rules for `evaluate`

An ordinary value

$$\frac{\mathcal{E}[M] = Ok\ V \quad M \not\equiv V}{\{\mathbf{E}[\mathbf{evaluate}\ M]\}_t \rightarrow \{\mathbf{E}[\mathbf{return}\ V]\}_t}$$

$$\frac{\mathcal{E}[M] = Bad\ S \quad e \in S}{\{\mathbf{E}[\mathbf{evaluate}\ M]\}_t \rightarrow \{\mathbf{E}[\mathbf{ioError}\ e]\}_t}$$

An exceptional value

Watch out!



$$\mathcal{E}[M] = V \quad M \not\equiv V$$
$$\frac{\{E[M]\}}{\{E[V]\}} \rightarrow \{E[V]\}$$

We've just changed
what values look
like!

But what if M
evaluates to (Bad S)???

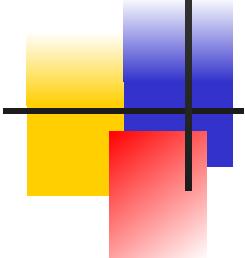
Watch out!

Ordinary value

Exceptional
value

$$\frac{\mathcal{E}[M] = Ok \ V \ M \not\equiv V}{\{\mathbf{E}[M]\}_t \rightarrow \{\mathbf{E}[V]\}_t} (FUN1)$$

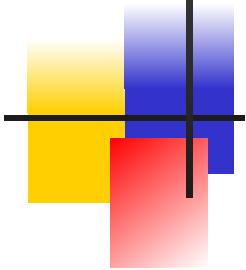
$$\frac{\mathcal{E}[M] = Bad \ S \ e \in S}{\{\mathbf{E}[M]\}_t \rightarrow \{\mathbf{E}[\text{ioError } e]\}_t} (FUN2)$$

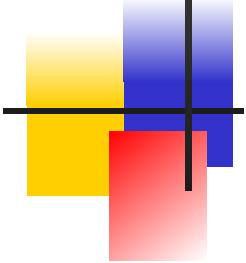


Imprecision exceptions

- A decent treatment of exceptions in purely-functional code
- Quite a lot more to say (see PLDI'99 paper)
 - No transformations lost!
- Good for disaster recovery, poor for extra return values

Asynchronous exceptions





Asynchronous exceptions

A flexible form of asynchronous exception:

```
throw :: Exception -> IO a  
throwTo :: ThreadId -> Exception ->  
          IO a
```



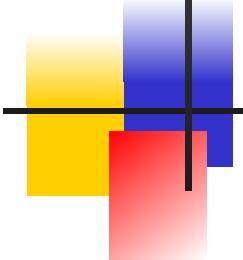
Timeouts

Fork a thread that sleeps and then throws an exception to its parent

```
timeout :: Int -> IO a -> IO (Maybe a)
timeout n a
= do { t <- myThreadId ;
      s <- forkIO (do { sleep n ;
                         throwTo t TimeOut }) ;
      catch (do { r <- a ;
                  throwTo s Kill ;
                  return (Just r) } ) ;
             (\ex -> Nothing)
      }
```

Do the action, and then kill the timeout

The timeout won!



Semantics

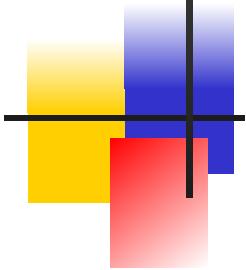
Add a rule for `throwTo`

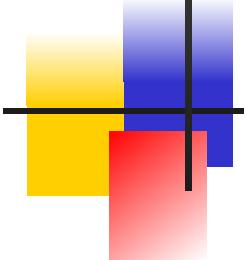
Make sure we
replace the
innermost
“current action”

$$\frac{M \neq (N_1 >= N_2) \quad M \neq (\text{catch } N_1 N_2)}{\{\text{E}_1[\text{throwTo } t \ e]\}_s \mid \{\text{E}_2[M]\}_t \rightarrow \{\text{E}_1[\text{return } ()]\}_s \mid \{\text{E}_2[\text{ioError } e]\}_t}$$

Replace “current
action” in target
thread with `ioError`

What have we achieved?





My motivation

Functional programming is SO much fun.

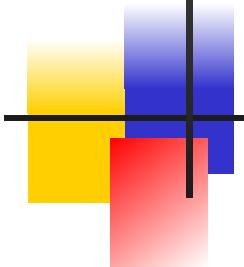
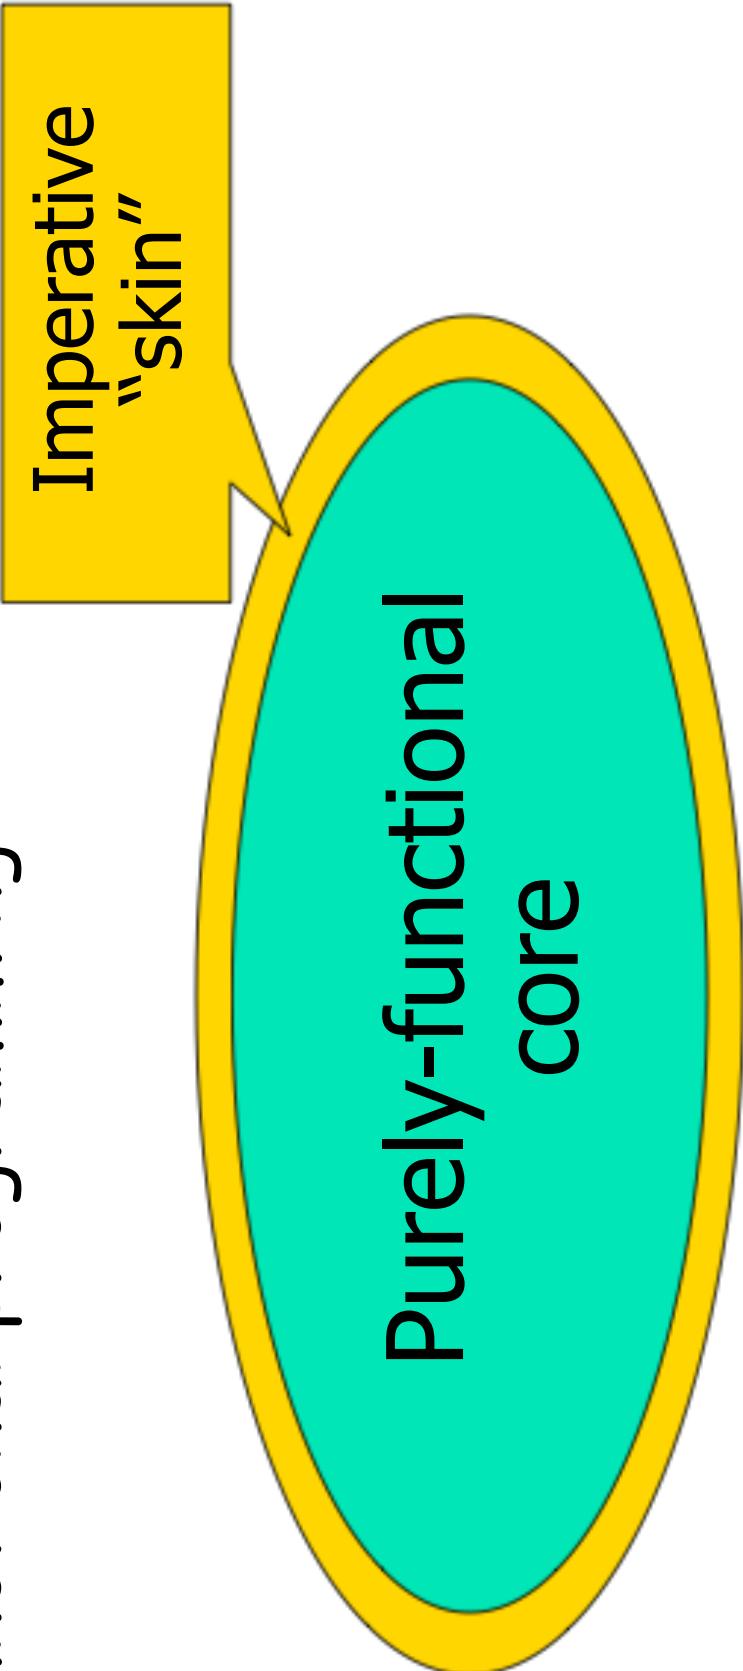
Plan of attack

- Find an application
- Try to write it in Haskell
- Fail
- Figure out how to fix Haskell
- Abstract key ideas, write a paper
- Repeat from (2)



What have we achieved?

- The ability to mix imperative and purely-functional programming



What have we achieved?

- ...without ruining either
 - All laws of pure functional programming remain unconditionally true, even of actions

e.g. $\text{let } x = e \text{ in } \dots x \dots x \dots$

=

$\dots e \dots e \dots$

What we have not achieved

- Imperative programming is no easier than it always was

e.g. `do { ...; x <- f 1; y <- f 2; ...}`

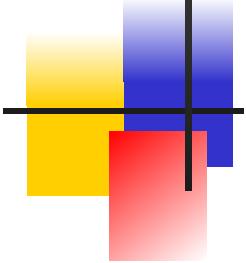
`?=?`

`do { ...; y <- f 2; x <- f 1; ... }`

...but there's less of it!

...and actions are first-class values





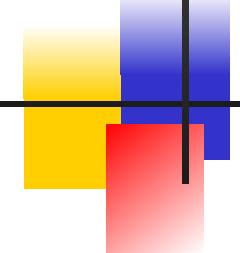
Not covered in the lectures

...But in the notes

Asynchronous exceptions

Foreign language interfacing





What next?

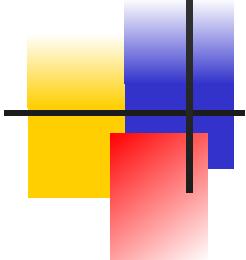
- Write applications

- Real reasoning about monadic Haskell programs; proving theorems

- Alternative semantic models (trace semantics)

- More refined monads (the IO monad is a giant sin-bin at the moment)

What next?



<http://research.microsoft.com/~simonpj>
<http://haskell.org>

Have lots
more fun